

### Task1

mysql> show databases;

| Database           |
|--------------------|
| information_schema |
| metastore          |
| mysql              |
| ngdb               |
| oozie              |
| performance_schema |
| simplidb           |
| sqooptest          |
| sys                |

9 rows in set (0.00 sec)

mysql> select \* from simplidb.Person;

| person_id | lname | fname  | area        | city          |
|-----------|-------|--------|-------------|---------------|
| 1         | Shyam | Ram    | Patna       | Bihar         |
| 2         | Tanya | Priya  | Whitefiled  | Bangalore     |
| 3         | James | Brown  | New York    | United States |
| 4         | Jhon  | Miller | Los Angeles | United States |
| 789       | a     | b      | c           | d             |

5 rows in set (0.00 sec)

mysql> show tables;

ERROR 1046 (3D000): No database selected

mysql> use simplidb;

Reading table information for completion of table and column names

You can turn off this feature to get a quicker startup with -A

Database changed

mysql> show tables;

| Tables_in_simplidb |
|--------------------|
| Person             |

1 row in set (0.00 sec)

mysql> create table student(id INT,fname varchar(50),lname varchar(50),area varchar(50),city varchar(50));

Query OK, 0 rows affected (0.02 sec)

mysql> describe student;

| Field | Type        | Null | Key | Default | Extra |
|-------|-------------|------|-----|---------|-------|
| id    | int(11)     | YES  |     | NULL    |       |
| fname | varchar(50) | YES  |     | NULL    |       |
| lname | varchar(50) | YES  |     | NULL    |       |

```
| area | varchar(50) | YES | | NULL | |
| city | varchar(50) | YES | | NULL | |
+-----+-----+-----+-----+-----+-----+
5 rows in set (0.00 sec)
```

```
[acadgild@localhost ~]$ sqoop import --connect
jdbc:mysql://localhost/simplidb --table Person --username
root -P --target-dir /sqooptesting2 -m 1
```

```
18/06/26 14:09:39 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
18/06/26 14:09:44 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/06/26 14:09:44 INFO tool.CodeGenTool: Beginning code generation
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/
slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/
slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-
impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple\_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
Tue Jun 26 14:09:45 IST 2018 WARN: Establishing SSL connection without server's identity verification
is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection
must be established by default if explicit option isn't set. For compliance with existing applications not
using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL
by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.
18/06/26 14:09:45 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person`
AS t LIMIT 1
18/06/26 14:09:45 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person`
AS t LIMIT 1
18/06/26 14:09:45 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/
hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadgild/compile/cd94a705db62cc92437fee228baa5448/Person.java uses or
overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/06/26 14:09:48 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/
cd94a705db62cc92437fee228baa5448/Person.jar
18/06/26 14:09:48 WARN manager.MySQLManager: It looks like you are importing from mysql.
18/06/26 14:09:48 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
18/06/26 14:09:48 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.
18/06/26 14:09:48 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull
(mysql)
18/06/26 14:09:48 INFO mapreduce.ImportJobBase: Beginning import of Person
18/06/26 14:09:48 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
18/06/26 14:09:48 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use
mapreduce.job.jar
18/06/26 14:09:49 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use
mapreduce.job.maps
18/06/26 14:09:49 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
Tue Jun 26 14:09:50 IST 2018 WARN: Establishing SSL connection without server's identity verification
is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection
```

must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.

18/06/26 14:09:51 INFO db.DBInputFormat: Using read committed transaction isolation

18/06/26 14:09:51 INFO mapreduce.JobSubmitter: number of splits:1

18/06/26 14:09:51 INFO mapreduce.JobSubmitter: Submitting tokens for job:  
job\_1530000379560\_0002

18/06/26 14:09:51 INFO impl.YarnClientImpl: Submitted application application\_1530000379560\_0002

18/06/26 14:09:51 INFO mapreduce.Job: The url to track the job: [http://localhost:8088/proxy/application\\_1530000379560\\_0002/](http://localhost:8088/proxy/application_1530000379560_0002/)

18/06/26 14:09:51 INFO mapreduce.Job: Running job: job\_1530000379560\_0002

18/06/26 14:09:59 INFO mapreduce.Job: Job job\_1530000379560\_0002 running in uber mode : false

18/06/26 14:09:59 INFO mapreduce.Job: map 0% reduce 0%

18/06/26 14:10:08 INFO mapreduce.Job: map 100% reduce 0%

18/06/26 14:10:08 INFO mapreduce.Job: Job job\_1530000379560\_0002 completed successfully

18/06/26 14:10:08 INFO mapreduce.Job: Counters: 30

File System Counters

- FILE: Number of bytes read=0
- FILE: Number of bytes written=127868
- FILE: Number of read operations=0
- FILE: Number of large read operations=0
- FILE: Number of write operations=0
- HDFS: Number of bytes read=87
- HDFS: Number of bytes written=150
- HDFS: Number of read operations=4
- HDFS: Number of large read operations=0
- HDFS: Number of write operations=2

Job Counters

- Launched map tasks=1
- Other local map tasks=1
- Total time spent by all maps in occupied slots (ms)=4943
- Total time spent by all reduces in occupied slots (ms)=0
- Total time spent by all map tasks (ms)=4943
- Total vcore-milliseconds taken by all map tasks=4943
- Total megabyte-milliseconds taken by all map tasks=5061632

Map-Reduce Framework

- Map input records=5
- Map output records=5
- Input split bytes=87
- Spilled Records=0
- Failed Shuffles=0
- Merged Map outputs=0
- GC time elapsed (ms)=47
- CPU time spent (ms)=1240
- Physical memory (bytes) snapshot=144351232
- Virtual memory (bytes) snapshot=2063667200
- Total committed heap usage (bytes)=127795200

File Input Format Counters

- Bytes Read=0

File Output Format Counters

- Bytes Written=150

18/06/26 14:10:08 INFO mapreduce.ImportJobBase: Transferred 150 bytes in 19.1228 seconds (7.844 bytes/sec)

18/06/26 14:10:08 INFO mapreduce.ImportJobBase: Retrieved 5 records.

You have new mail in [var/spool/mail/acadgild](mailto:var/spool/mail/acadgild)

[acadgild@localhost ~]\$ hdfs dfs -cat [/sqoop-testing2/part-m-0000018/06/26](#) 14:10:28 WARN

util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java

classes where applicable

- 1, Shyam, Ram, Patna, Bihar
  - 2, Tanya, Priya, Whitefield, Bangalore
  - 3, James, Brown, New York, United States
  - 4, Jhon, Miller, Los Angeles, United States
- 789, a, b, c, d

```
[acadgild@localhost ~]$ sqoop export --connect
jdbc:mysql://localhost/simplidb --username root --password
Root@123 --table student --export-dir /sqooptesting2/part-
m-00000 -m 1
```

```
18/06/26 14:29:19 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
18/06/26 14:29:19 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure.
Consider using -P instead.
18/06/26 14:29:19 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/06/26 14:29:19 INFO tool.CodeGenTool: Beginning code generation
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/
slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/
slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-
impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple\_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
Tue Jun 26 14:29:20 IST 2018 WARN: Establishing SSL connection without server's identity verification
is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection
must be established by default if explicit option isn't set. For compliance with existing applications not
using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL
by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.
18/06/26 14:29:21 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `student`
AS t LIMIT 1
18/06/26 14:29:21 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `student`
AS t LIMIT 1
18/06/26 14:29:21 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/
hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadgild/compile/afa70f4108efcc41706e89e8507a259d/student.java uses or
overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/06/26 14:29:23 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/
afa70f4108efcc41706e89e8507a259d/student.jar
18/06/26 14:29:23 INFO mapreduce.ExportJobBase: Beginning export of student
18/06/26 14:29:23 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
18/06/26 14:29:23 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use
mapreduce.job.jar
18/06/26 14:29:24 INFO Configuration.deprecation: mapred.reduce.tasks.speculative.execution is
deprecated. Instead, use mapreduce.reduce.speculative
18/06/26 14:29:24 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is
deprecated. Instead, use mapreduce.map.speculative
```

18/06/26 14:29:24 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use  
 mapreduce.job.maps  
 18/06/26 14:29:24 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032  
 18/06/26 14:29:26 INFO input.FileInputFormat: Total input paths to process : 1  
 18/06/26 14:29:26 INFO input.FileInputFormat: Total input paths to process : 1  
 18/06/26 14:29:26 INFO mapreduce.JobSubmitter: number of splits:1  
 18/06/26 14:29:26 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is  
 deprecated. Instead, use mapreduce.map.speculative  
 18/06/26 14:29:26 INFO mapreduce.JobSubmitter: Submitting tokens for job:  
 job\_1530000379560\_0003  
 18/06/26 14:29:26 INFO impl.YarnClientImpl: Submitted application application\_1530000379560\_0003  
 18/06/26 14:29:26 INFO mapreduce.Job: The url to track the job: [http://localhost:8088/proxy/  
 application\\_1530000379560\\_0003/](http://localhost:8088/proxy/application_1530000379560_0003/)  
 18/06/26 14:29:26 INFO mapreduce.Job: Running job: job\_1530000379560\_0003  
 18/06/26 14:29:34 INFO mapreduce.Job: Job job\_1530000379560\_0003 running in uber mode : false  
 18/06/26 14:29:34 INFO mapreduce.Job: map 0% reduce 0%  
 18/06/26 14:29:40 INFO mapreduce.Job: map 100% reduce 0%  
 18/06/26 14:29:41 INFO mapreduce.Job: Job job\_1530000379560\_0003 completed successfully  
 18/06/26 14:29:41 INFO mapreduce.Job: Counters: 30  
     File System Counters  
         FILE: Number of bytes read=0  
         FILE: Number of bytes written=127555  
         FILE: Number of read operations=0  
         FILE: Number of large read operations=0  
         FILE: Number of write operations=0  
         HDFS: Number of bytes read=282  
         HDFS: Number of bytes written=0  
         HDFS: Number of read operations=4  
         HDFS: Number of large read operations=0  
         HDFS: Number of write operations=0  
     Job Counters  
         Launched map tasks=1  
         Data-local map tasks=1  
         Total time spent by all maps in occupied slots (ms)=3939  
         Total time spent by all reduces in occupied slots (ms)=0  
         Total time spent by all map tasks (ms)=3939  
         Total vcore-milliseconds taken by all map tasks=3939  
         Total megabyte-milliseconds taken by all map tasks=4033536  
     Map-Reduce Framework  
         Map input records=5  
         Map output records=5  
         Input split bytes=129  
         Spilled Records=0  
         Failed Shuffles=0  
         Merged Map outputs=0  
         GC time elapsed (ms)=43  
         CPU time spent (ms)=930  
         Physical memory (bytes) snapshot=141643776  
         Virtual memory (bytes) snapshot=2061561856  
         Total committed heap usage (bytes)=127795200  
     File Input Format Counters  
         Bytes Read=0  
     File Output Format Counters  
         Bytes Written=0  
 18/06/26 14:29:41 INFO mapreduce.ExportJobBase: Transferred 282 bytes in 17.059 seconds (16.5308  
 bytes/sec)  
 18/06/26 14:29:41 INFO mapreduce.ExportJobBase: Exported 5 records.

```
mysql> show tables;
+-----+
| Tables_in_simplidb |
+-----+
| Person              |
| student              |
+-----+
2 rows in set (0.00 sec)
```

```
mysql> select * from simplidb.student;
+-----+-----+-----+-----+-----+
| id | fname | lname | area | city |
+-----+-----+-----+-----+
| 1 | Shyam | Ram | Patna | Bihar |
| 2 | Tanya | Priya | Whitefield | Bangalore |
| 3 | James | Brown | New York | United States |
| 4 | Jhon | Miller | Los Angeles | United States |
| 789 | a | b | c | d |
+-----+-----+-----+-----+
5 rows in set (0.00 sec)
```

## Task2

```
[acadgild@localhost ~]$ sqoop import --connect jdbc:mysql://localhost/simplidb --username root -P --
query 'select * from student where id = 3 and $CONDITIONS' --target-dir /sqooptesting5 -m 1
18/06/26 15:22:06 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
18/06/26 15:22:17 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/06/26 15:22:17 INFO tool.CodeGenTool: Beginning code generation
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/
slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/
slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-
impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple\_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
Tue Jun 26 15:22:17 IST 2018 WARN: Establishing SSL connection without server's identity verification
is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection
must be established by default if explicit option isn't set. For compliance with existing applications not
using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL
by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.
18/06/26 15:22:18 INFO manager.SqlManager: Executing SQL statement: select * from student where
id = 3 and (1 = 0)
18/06/26 15:22:18 INFO manager.SqlManager: Executing SQL statement: select * from student where
id = 3 and (1 = 0)
18/06/26 15:22:18 INFO manager.SqlManager: Executing SQL statement: select * from student where
id = 3 and (1 = 0)
18/06/26 15:22:18 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadgild/compile/cbaecc782902d13f3421c45baf88f23d/QueryResult.java uses or
overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/06/26 15:22:20 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/cbaecc782902d13f3421c45baf88f23d/QueryResult.jar
18/06/26 15:22:20 INFO mapreduce.ImportJobBase: Beginning query import.
```



18/06/26 15:22:20 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable  
18/06/26 15:22:21 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar  
18/06/26 15:22:21 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps  
18/06/26 15:22:21 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032  
Tue Jun 26 15:22:23 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.  
18/06/26 15:22:23 INFO db.DBInputFormat: Using read committed transaction isolation  
18/06/26 15:22:23 INFO mapreduce.JobSubmitter: number of splits:1  
18/06/26 15:22:23 INFO mapreduce.JobSubmitter: Submitting tokens for job: job\_1530000379560\_0006  
18/06/26 15:22:24 INFO impl.YarnClientImpl: Submitted application application\_1530000379560\_0006  
18/06/26 15:22:24 INFO mapreduce.Job: The url to track the job: [http://localhost:8088/proxy/application\\_1530000379560\\_0006/](http://localhost:8088/proxy/application_1530000379560_0006/)  
18/06/26 15:22:24 INFO mapreduce.Job: Running job: job\_1530000379560\_0006  
18/06/26 15:22:32 INFO mapreduce.Job: Job job\_1530000379560\_0006 running in uber mode : false  
18/06/26 15:22:32 INFO mapreduce.Job: map 0% reduce 0%  
18/06/26 15:22:38 INFO mapreduce.Job: map 100% reduce 0%  
18/06/26 15:22:38 INFO mapreduce.Job: Job job\_1530000379560\_0006 completed successfully  
18/06/26 15:22:38 INFO mapreduce.Job: Counters: 30

#### File System Counters

FILE: Number of bytes read=0  
FILE: Number of bytes written=127556  
FILE: Number of read operations=0  
FILE: Number of large read operations=0  
FILE: Number of write operations=0  
HDFS: Number of bytes read=87  
HDFS: Number of bytes written=37  
HDFS: Number of read operations=4  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=2

#### Job Counters

Launched map tasks=1  
Other local map tasks=1  
Total time spent by all maps in occupied slots (ms)=3923  
Total time spent by all reduces in occupied slots (ms)=0  
Total time spent by all map tasks (ms)=3923  
Total vcore-milliseconds taken by all map tasks=3923  
Total megabyte-milliseconds taken by all map tasks=4017152

#### Map-Reduce Framework

Map input records=1  
Map output records=1  
Input split bytes=87  
Spilled Records=0  
Failed Shuffles=0  
Merged Map outputs=0  
GC time elapsed (ms)=47  
CPU time spent (ms)=1150  
Physical memory (bytes) snapshot=144080896  
Virtual memory (bytes) snapshot=2063667200  
Total committed heap usage (bytes)=127795200

#### File Input Format Counters

Bytes Read=0  
File Output Format Counters  
Bytes Written=37

18/06/26 15:22:38 INFO mapreduce.ImportJobBase: Transferred 37 bytes in 17.0951 seconds (2.1644 bytes/sec)

18/06/26 15:22:38 INFO mapreduce.ImportJobBase: Retrieved 1 records.

You have new mail in </var/spool/mail/acadgild>

[acadgild@localhost ~]\$ hdfs dfs -cat </sqoop/testing5/part-m-0000018/06/26> 15:22:51 WARN  
util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java  
classes where applicable

3,James,Brown,New York,United States

using the export command we can transfer this file to RDBMS file system (as done for [task1](#))